




Introduction to MongoDB

Database compared



What is MongoDB ?

- Scalable High-Performance Open-source, Document-orientated database.
- Built for Speed
- Rich Document based queries for **Easy readability**.
- Full Index Support for **High Performance**.
- Replication and Failover for **High Availability**.
- Auto Sharding for **Easy Scalability**.


Why use MongoDB?

- SQL was invented in the 70's to store **data**.
- MongoDB stores **documents (or) objects**.
- Now-a-days, everyone works with **objects** (Python/Ruby/Java/etc.)
- And we need Databases to persist our **objects**. Then why not store **objects** directly ?
- Embedded documents and arrays reduce need for joins. **No Joins** and No-multi document **transactions**.



What is MongoDB great for?

- RDBMS replacement for **Web Applications**.
- **Semi-structured** Content Management.
- **Real-time** Analytics & High-Speed Logging.
- Caching and **High Scalability**



Not great for?

- Highly **Transactional** Applications.
- Problems requiring **SQL**.

Impedance Mismatch

```
// your application code
class Foo { int x; string
[] tags;}
```

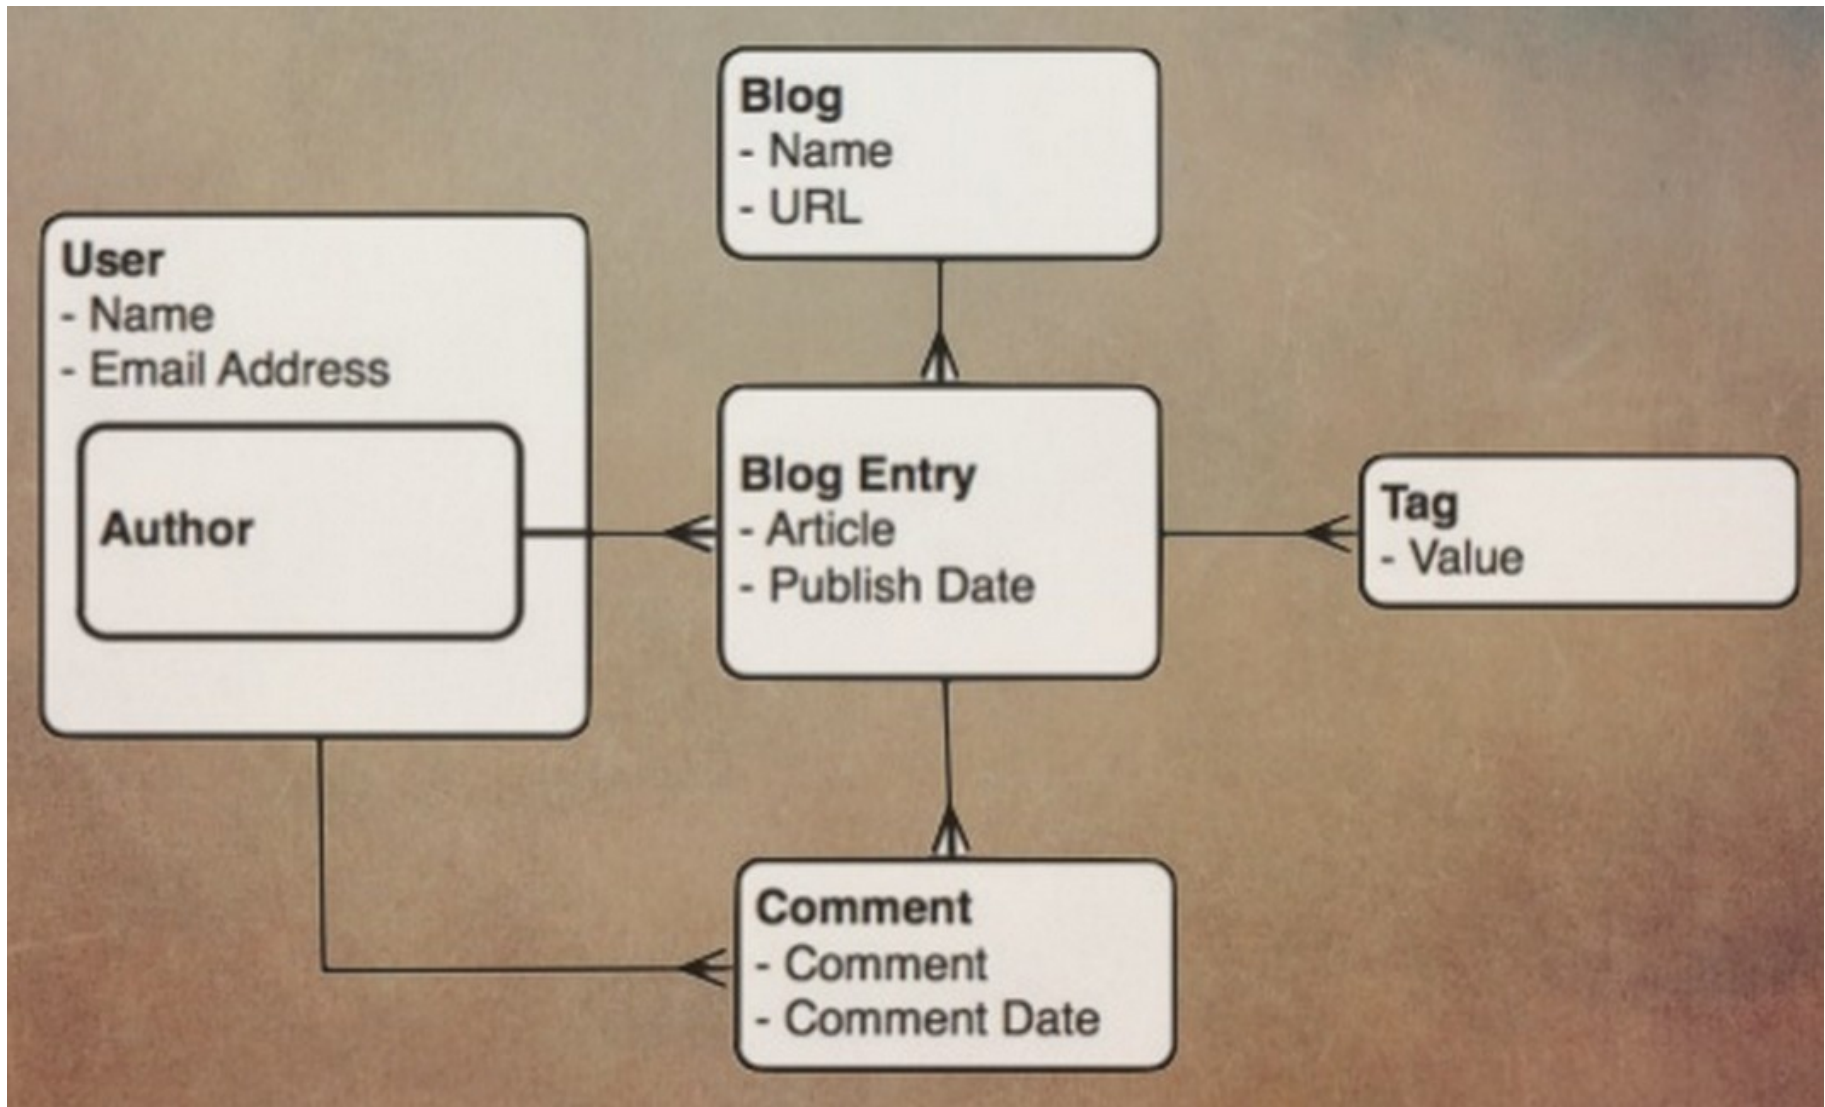
x	nam		
1	Ab		
2	xy		
	z	tagId	
		3	1
		3	2
		4	2
tagId	tag		
3	red		
3	blu		
4	e		

No Impedance Mismatch

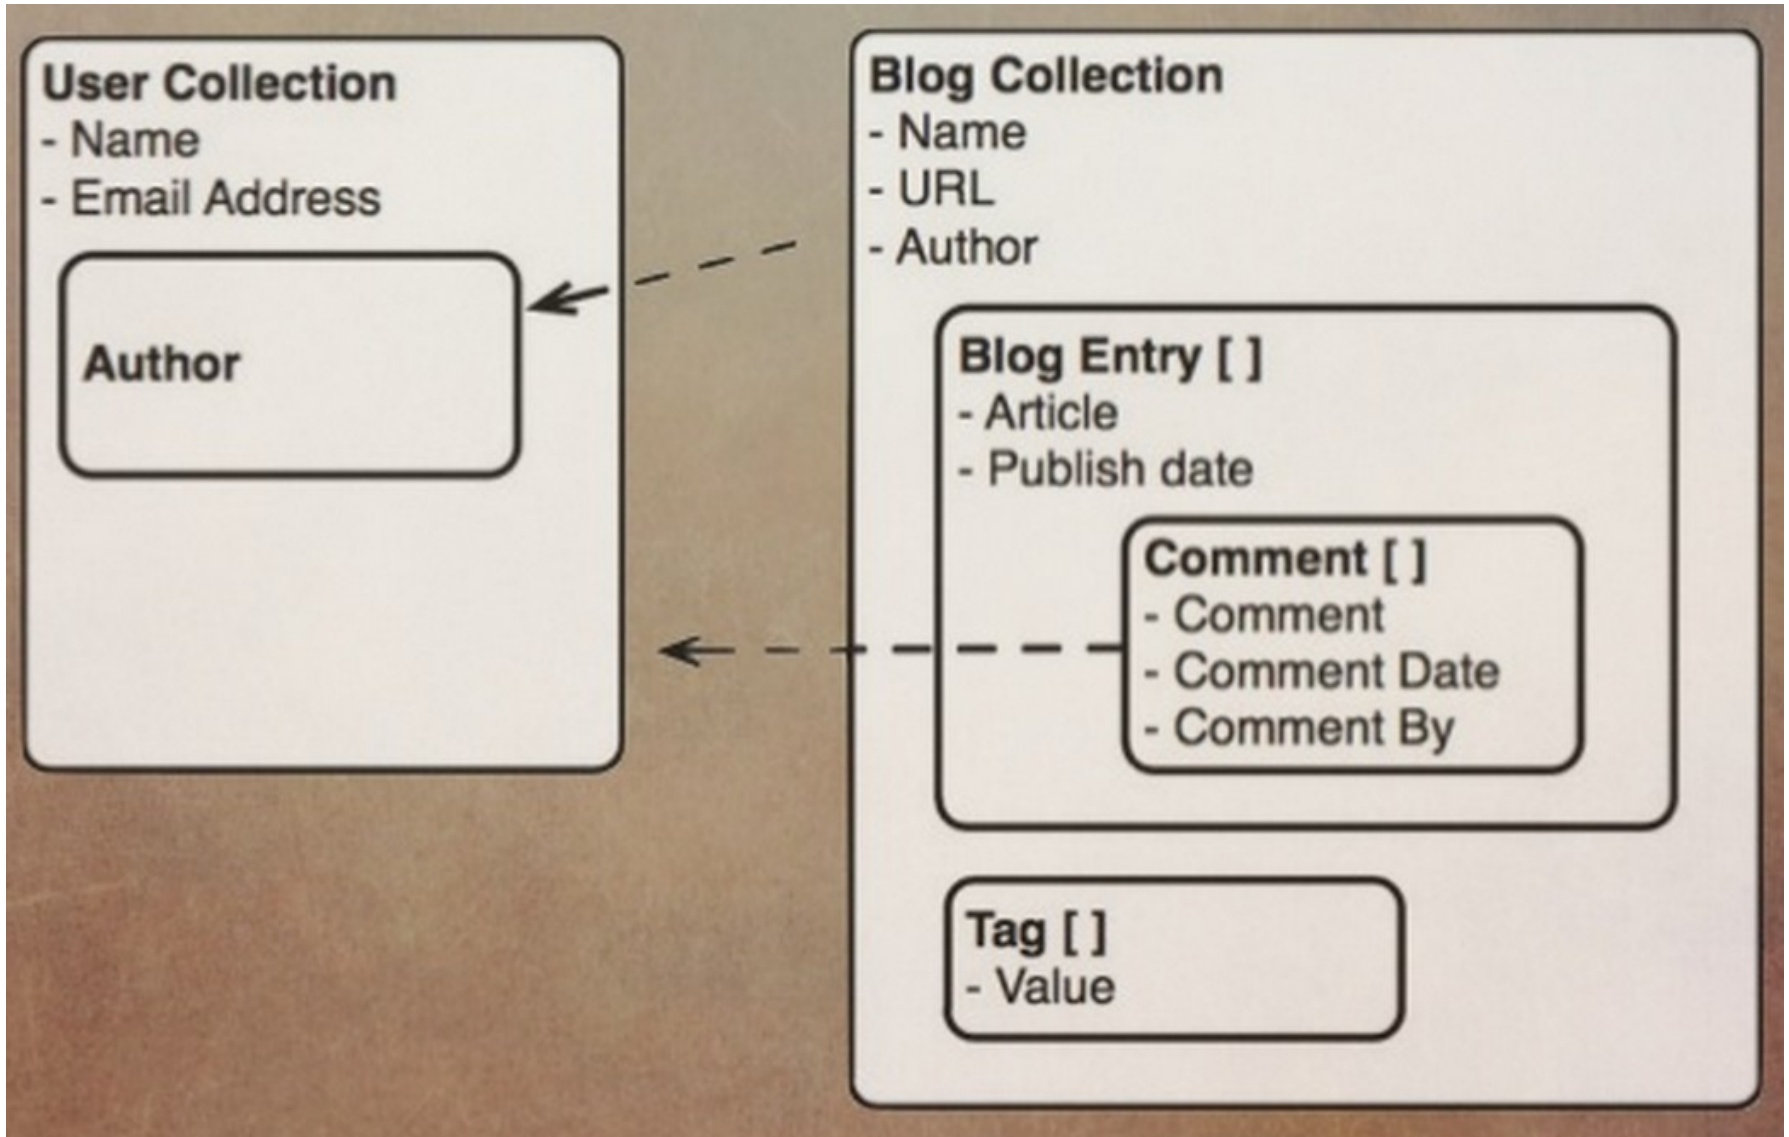
```
// your application code
class Foo { int x; string [] tags;}

// mongo document for Foo
{ x: 1, tags: ['abc','xyz'] }
```


Blog in relational DB



Blog post structure in document DB



Blog post in JSON DB

```
{  _id : ObjectId("4c4ba5c0672c685e5e8aabf3"),
  author : "steve",
  date : "Sat Apr 24 2013 19:47:11",
  text : "About MongoDB...",
  tags : [ "tech", "databases" ],
  comments : [
    {
      author : "Fred",
      date : "Sat Apr 25 2013 20:51:03 GMT-0700",
      text : "Best Post Ever!"
    }
  ]
}
```

When I say
Database



Think
Database

- Made up of Multiple **Collections**.
- Created **on-the-fly** when referenced for the first time.

When I say
Collection



Think
Table

- Schema-less, and contains **Documents**.
- **Indexable** by one/more keys.
- Created **on-the-fly** when referenced for the first time.
- **Capped Collections**: Fixed size, older records get dropped after reaching the limit.

When I say
Document



Think
Record/Row

- Stored in a **Collection**.
- Have **_id** key – works like Primary keys in MySQL.
- Supported Relationships – **Embedded (or) References**.
- Document storage in **BSON** (Binary form of JSON).

Understanding the Document Model

```
var post = {  
  '_id': ObjectId('3432'),  
  'author': ObjectId('2311'),  
  'title': 'Introduction to MongoDB',  
  'body': 'MongoDB is an open sources.. ',  
  'timestamp': Date('01-04-12'),  
  'tags': ['MongoDB', 'NoSQL'],  
  'comments': [{ 'author': ObjectId('5331'),  
                  'date': Date('02-04-12'),  
                  'text': 'Did you see.. ',  
                  'upvotes': 7 } ]  
}  
  
> db.posts.insert(post);
```

The Problem

- **You**

```
db.foo.find({ x:  
10 })
```

- **The server does :**

```
for each doc d in 'foo'{  
    if ( d.x == 10 ){  
        return d  
    }  
}
```

Ouc
Reads ~~the~~ EVERY
document!

Document Storage

_id: ...
x:9

_id: ...
x:9

...

_id: ...
x:9

_id: ...
x:9

...

_id: ...
x:9

_id: ...
x:9

...

_id: ...
x:10

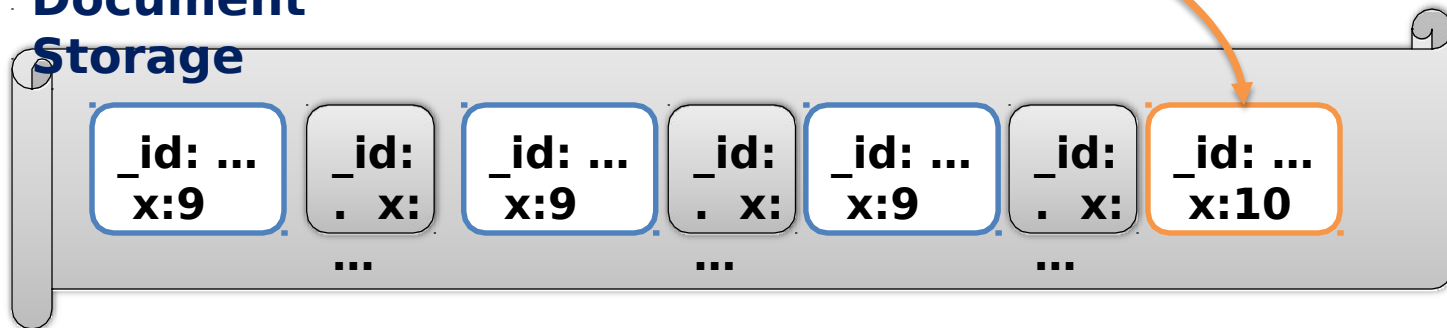
The Solution

Index - field 'x', collection

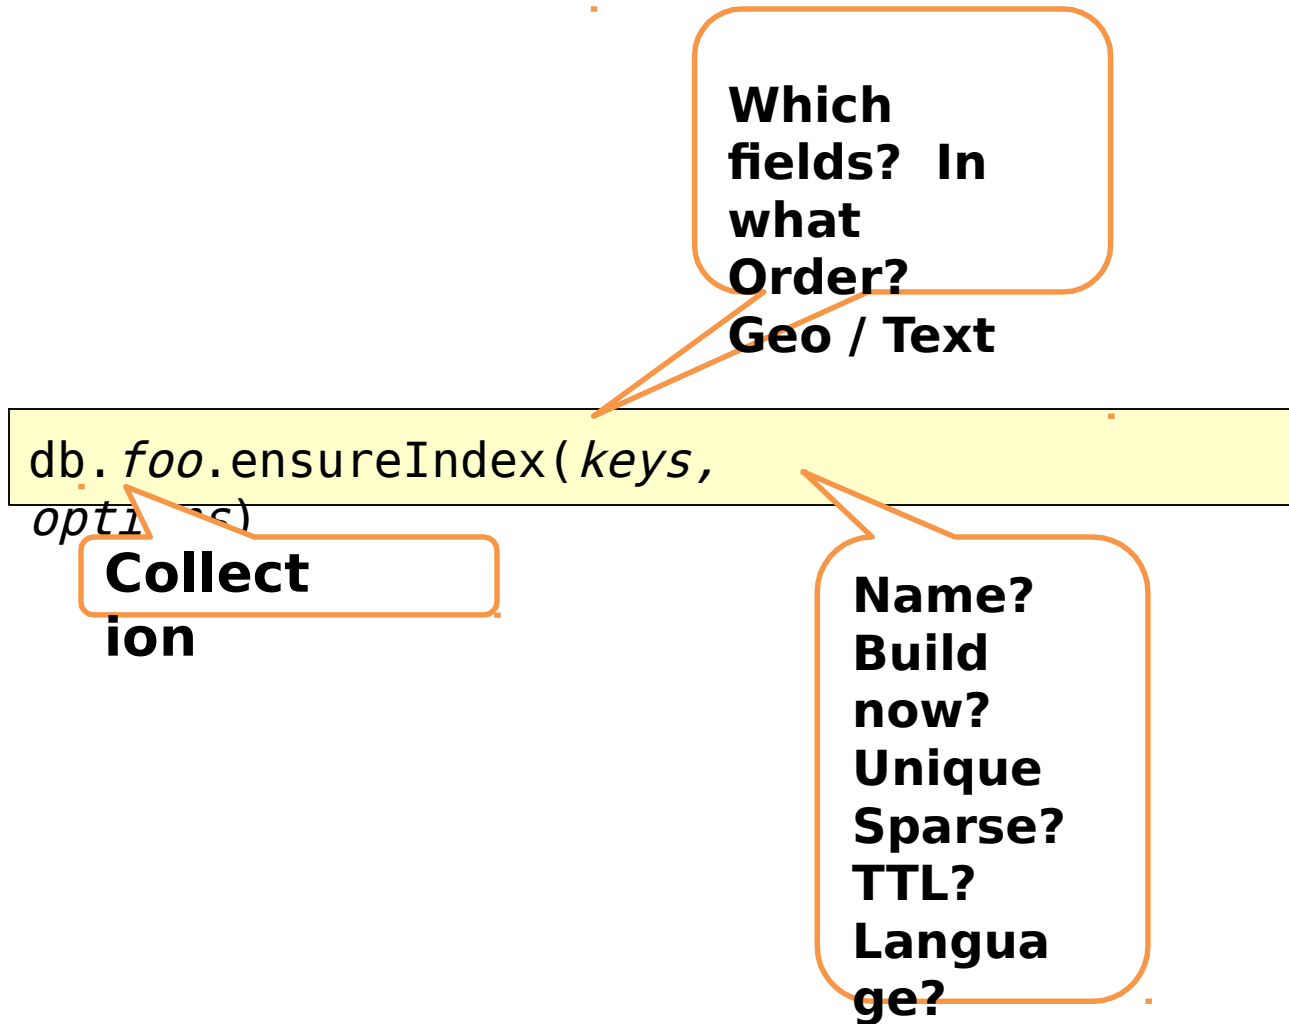
Value	Doc Pointers
9	[171, 819, 2309]
10	[4376]

```
db.foo.find({ x:10 })
```

**Document
Storage**



Create Index



Secondary Indexes

Create Index on any field in the document

// 1 means ascending, -1 means descending

```
> db.posts.ensureIndex({'author': 1});
```

//Index Nested Documents

```
> db.posts.ensureIndex('comments.author': 1);
```

// Index on tags

```
> db.posts.ensureIndex({'tags': 1});
```

// Geo-spatial Index

```
> db.posts.ensureIndex({'author.location': '2d'});
```

Find

// find posts which has 'MongoDB' tag.

```
> db.posts.find({tags: 'MongoDB'});
```

// find posts by author's comments.


```
> db.posts.find({'comments.author': 'Johnson'}).count();
```

// find posts written after 31st March.

```
> db.posts.find({'timestamp': {'$gte': Date('31-03-12')}});
```


// find posts written by authors around [22, 42]

```
> db.posts.find({'author.location': {'$near':[22, 42]}},  
  $gt, $lt, $gte, $lte, $ne, $all, $in, $nin...)
```



Find

**Which
fields?**



`db.foo.find(
 query,
 projection)`

**Which
documents?**

Find: projection

```
> db.posts.find({}, {title:1})
```

```
{ "_id" : ObjectId("5654381f37f63ffc4ebf1964"),  
  "title" : "NodeJS server" }  
{ "_id" : ObjectId("5654385c37f63ffc4ebf1965"),  
  "title" : "Introduction to MongoDB" }
```

Like

```
select title from posts
```

Empty projection like

```
select * from posts
```

Find

Find
d

- **Query criteria**
- **Single value field**
- **Array field**
- **Sub-document / dot notation**

Projecti
on

- **Filed inclusion and exclusion**

Curs
or

- **Sort**
- **Limit**
- **Skip**

Paging example

```
place1 = {  
    name : "10gen HQ",  
    address : "229 W 43rd St. 5th Floor",  
    city : "New York",  
    zip : "10036",  
    tags : [ "business", "awesome" ]  
}  
  
> db.places.insert(place1)
```

```
per_page = 10;  
page_num = 3;  
  
places = db.places  
    .find({ "city" : "new york" })  
    .sort({ "ts" : -1 })  
    .skip((page_num - 1) * per_page)  
    .limit(per_page);
```


Update: replace the document

```
> db.posts.update(  
  {"_id" : ObjectId("5654381f37f63ffc4ebf1964")},  
  {  
    title:"NodeJS server"  
  });
```

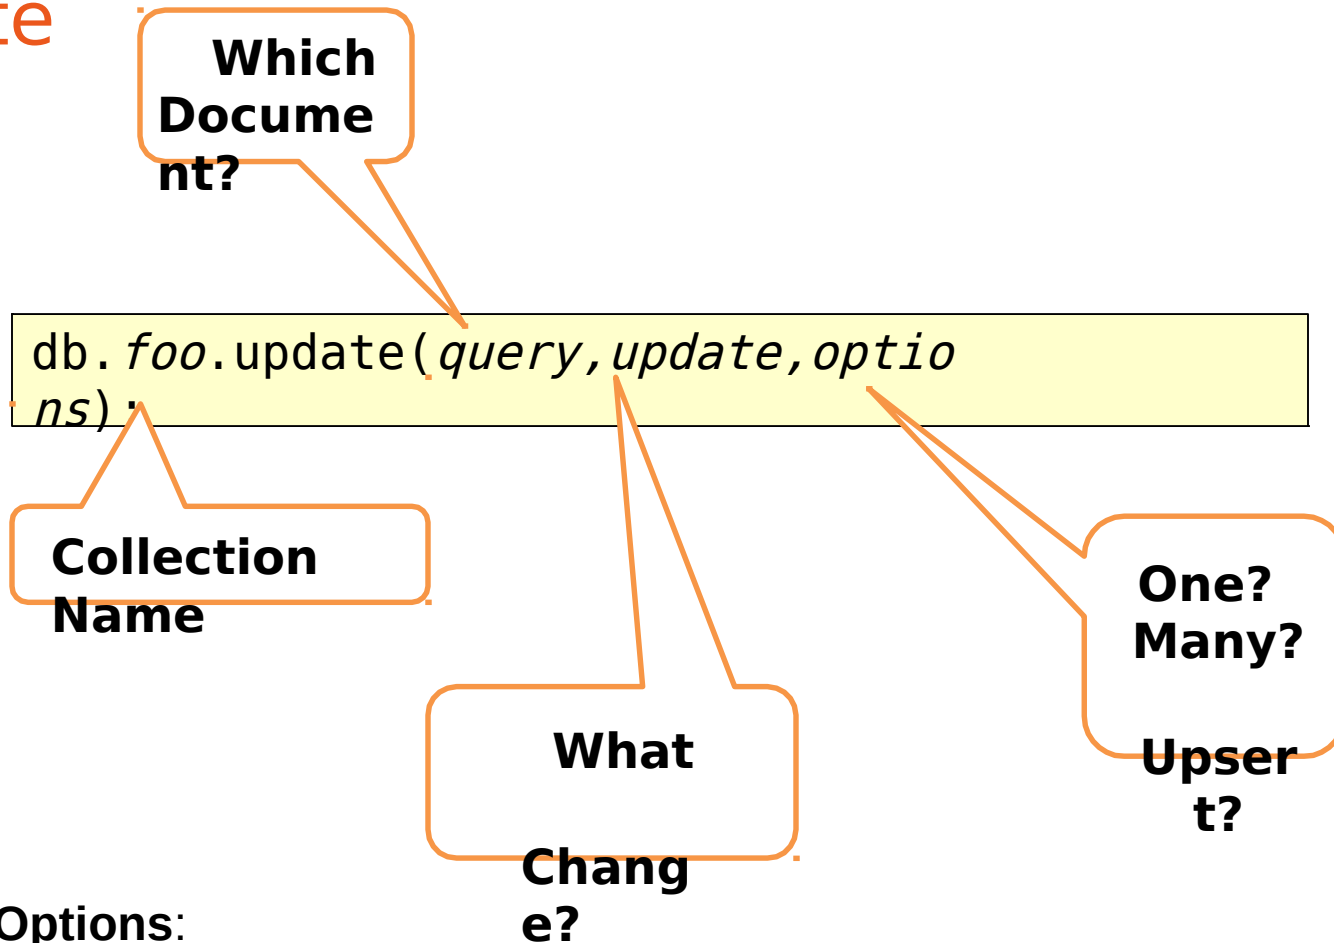
This will **replace** the document by {title:"NodeJS server"}

Update: change only the part of document

```
> db.posts.update(  
  {"_id" : ObjectId("5654381f37f63ffc4ebf1964")},  
  {  
    $addToSet: {tags:"JS"},  
    $set: {title:"NodeJS server"},  
    $unset: { comments: 1}  
  });
```

\$set, \$unset
\$push, \$pull, \$pop, \$addToSet
\$inc, \$decr, many more...

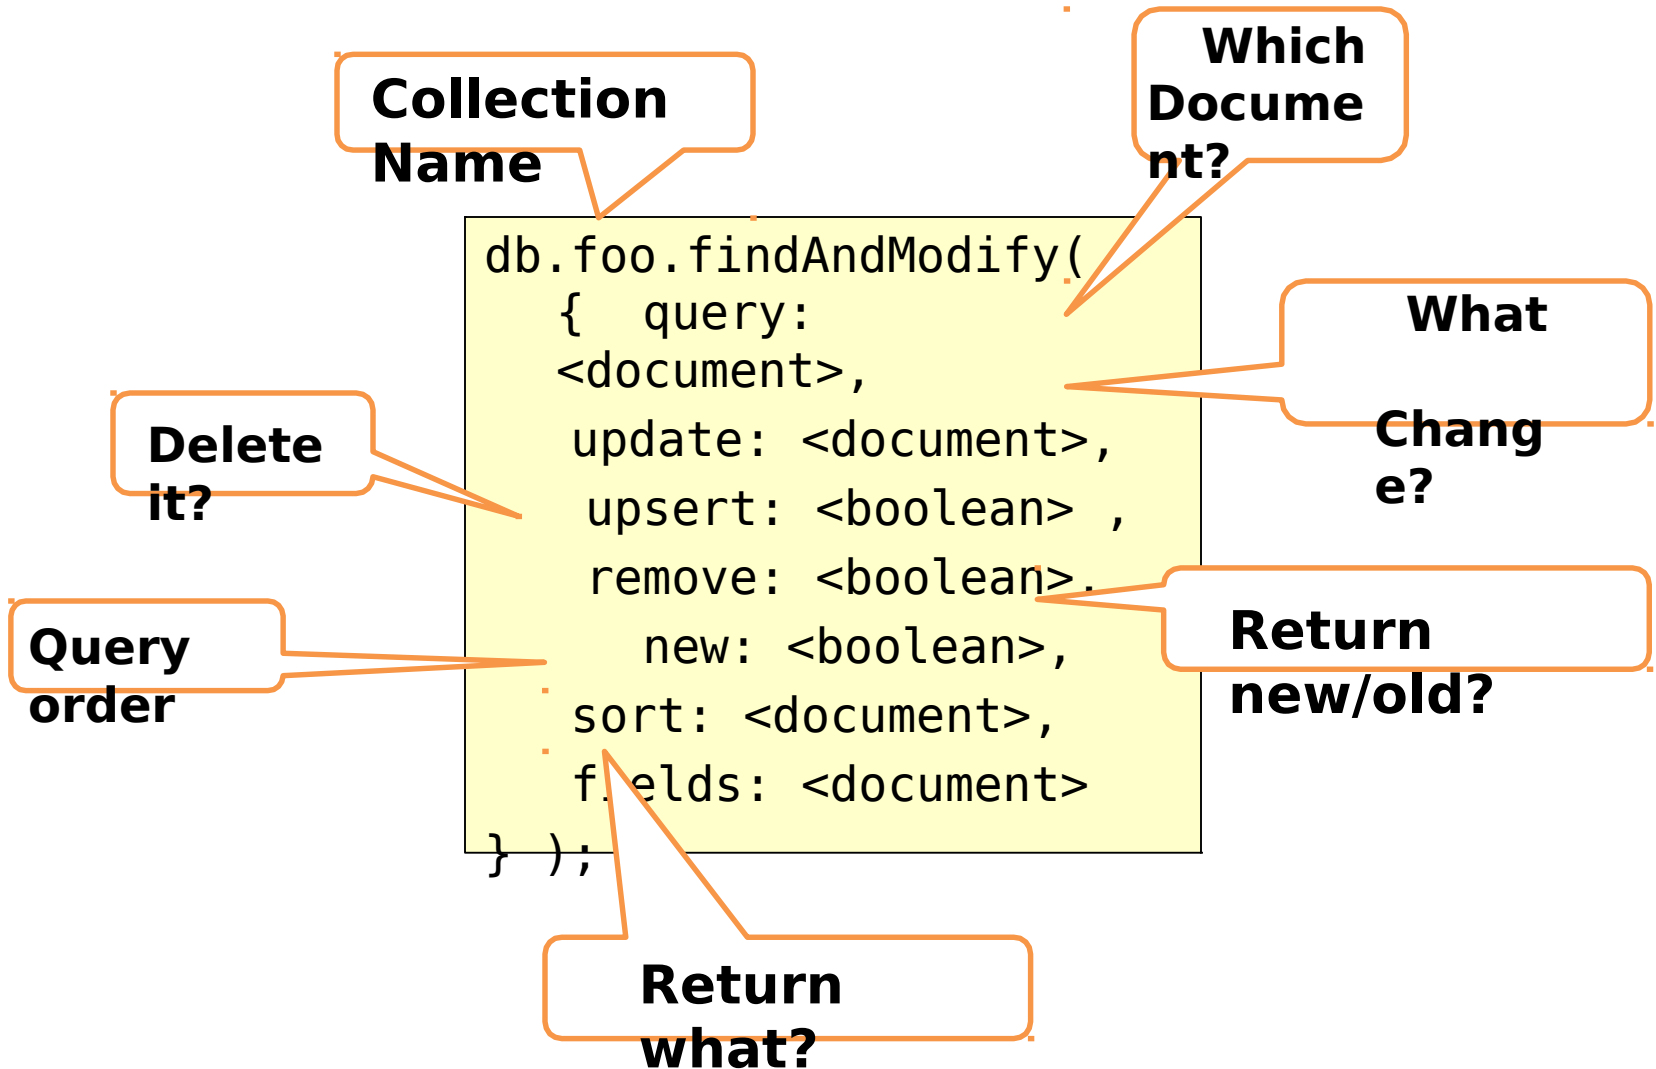
Update



Options:

- {multi: true}** – will change all found documents;
by default only first found will be updated
- {upsert: true}** – will insert document if it was not found

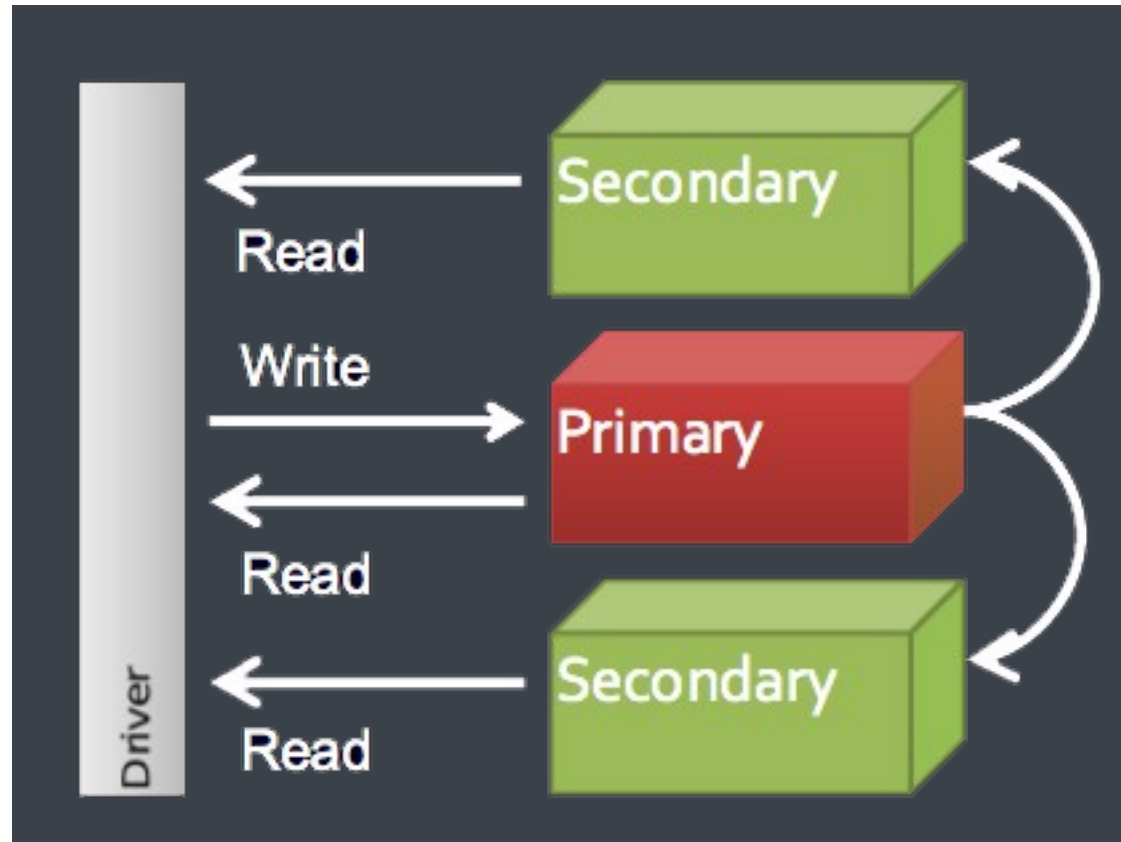
Find And Modify



Some Cool features

- Geo-spatial Indexes for Geo-spatial queries.
\$near, \$within_distance, Bound queries (circle, box)
- GridFS
Stores Large Binary Files.
- Map/Reduce
GROUP BY in SQL, map/reduce in MongoDB.

Replica Sets



Sharding

